# INKLOOM AI

Brand Governance Platform

ROUND 2 — THE BUILD

A functional MVP for checking brand consistency before content is published, explaining possible issues, and keeping human reviewers in control.

Project report · October 2026

# 1. Problem and solution

Organisations create content through multiple people and channels. Inconsistent terminology, risky promotional claims, vague link text, and tone mismatches can weaken trust. Inkloom AI is a lightweight brand-governance workspace that checks submitted text against configurable local rules and routes it to a human reviewer.

# 2. Target users

Content creators, marketing teams, brand managers, communications teams, and authorised reviewers who need a repeatable pre-publication check.

# 3. Completed Round 2 features

| Feature | What works in this MVP |
|---|---|
| Brand guidelines | Edit brand voice, required terminology, avoid-list terms, approved-claim notes, and accessibility guidance. |
| Content review | Submit titled content and choose a channel; local checks run on submission. |
| Explainable findings | Each flag includes severity, reason, and a suggested action. |
| Suggested draft | Conservative replacement of obvious avoid-list and absolute wording; review remains manual. |
| Approval workflow | Pending approval, approve, request revision, edit copy, and record reviewer notes. |
| Governance insights | Count repeated issue types across saved reviews. |
| Persistence | SQLite database stores guidelines and review history. |

# 4. Main user workflow

1. A manager configures the brand rules and reference notes.
2. A creator submits content and chooses a channel.
3. The system checks terminology, avoid-list phrases, absolute/numeric claims, vague links, and possible tone issues.
4. Findings explain why they were raised and what the reviewer can do.
5. A human edits the content and marks it Approved, Needs Revision, or Pending Approval.
6. The dashboard and insights page summarise saved reviews and repeated findings.

# 5. Technology and architecture

**Frontend:** HTML templates and responsive CSS. **Backend:** Python and Flask. **Database:** SQLite. **Analysis:** local, explainable rule checks using Python string and regular-expression matching. **External APIs:** none in this MVP.

The app initialises the SQLite database automatically on startup. The `/health` endpoint returns a small status response. No user content is sent to an external AI service.

# 6. Testing checklist

| Test | Expected result |
| --- | --- |
| Open home page | Dashboard loads and displays review counts. |
| Submit empty form | Validation asks for title and content. |
| Submit content with avoid-list wording | A finding is shown with reason and action. |
| Edit brand guidelines | Saved rules are used for subsequent reviews. |
| Approve or request revision | Status, content, note, and updated time are saved. |
| Open insights | Repeated findings and average heuristic score are displayed. |

# 7. Three major Round 3 improvements

**1. Evidence-linked claim verification.** Let reviewers attach approved source documents or URLs to claims and show the source beside each finding. Avoid declaring claims true without evidence.

**2. Image and document review.** Add image alt-text/accessibility checks and text extraction for uploaded DOCX/PDF files, with clear warnings when visual analysis is not supported.

**3. Team roles and audit trail.** Add creator/reviewer/admin permissions, a durable change log, reviewer identity, and timestamps for every status change.

# 8. Privacy, safety, and limitations

The current version runs locally and uses SQLite. It does not call an external AI model. Avoid-list matches and regex checks are heuristic and can produce false positives or miss context. Approved claims are stored as notes but are not semantically validated. The score is not a confidence probability, legal opinion, or guarantee of compliance. Human approval is required before publication.

# 9. Deployment and submission

Run locally using the steps in README.md, test the complete workflow, and then deploy to a hosting provider you control. Create a public source repository and verify the project URL and repository can be opened in a private/incognito browser without signing in. Do not submit localhost as the working project link.

AI/tooling disclosure: This project was scaffolded with AI assistance. The team should review, test, understand, and adapt the code, and follow the organiser's disclosure rules. The event states that only meaningful work completed during the official build window is judged; represent development history honestly.